



Module 8 - Storage & Networking

By Juan Medina

jmedina@collin.edu






Storage

Storage management looks a lot like a giant set of Lego blocks that you can put together in an infinite variety of configurations. Traditional hard disks remain the dominant medium for storage, but they're increasingly being joined by solid-state drives (SSDs) for performance-sensitive applications. On top of this hardware is a variety of software components like drivers, partitioning, RAID implementations, Logical Volume Managers, network storage and virtualization that mediate between the raw storage devices and the file-system hierarchy seen by users.

Disk Partitions



Storage devices are typically divided up into smaller chunks called partitions. Partitions allow you to organize a disk. The various partitions can be formatted with different file systems or used for different purposes. For example, one partition can contain user home directories while another can contain system data and logs.

Partitions are block devices in their own right.


A long listing of the `/dev/sd*` directory reveals its special file type as `b`, which stands for block device:

```

1 [jmedinar@localhost ~]$ ls -l /dev/sd*
2 brw-rw----. 1 root disk 8,  0 Dec 28 11:01 /dev/sda
3 brw-rw----. 1 root disk 8,  1 Dec 28 11:01 /dev/sda1
4 brw-rw----. 1 root disk 8,  2 Dec 28 11:01 /dev/sda2
5 brw-rw----. 1 root disk 8, 16 Dec 28 11:01 /dev/sdb
6 brw-rw----. 1 root disk 8, 17 Dec 28 11:01 /dev/sdb1
7 brw-rw----. 1 root disk 8, 32 Dec 28 11:01 /dev/sdc
8 brw-rw----. 1 root disk 8, 33 Dec 28 11:01 /dev/sdc1
9 brw-rw----. 1 root disk 8, 34 Dec 28 17:32 /dev/sdc2

```

Logical Volumes



Another way of organizing disks and partitions is with logical volume management (LVM). With LVM, one or more block devices can be aggregated into a storage pool called a volume group. Disk space in the volume group is then parceled out to one or more logical volumes, which are the functional equivalent of a partition residing on a physical disk.

File systems

Files on a Linux server are accessed through the file-system hierarchy, a single inverted tree of directories. Each file system is a storage device that has been formatted to store files.

In a sense, the Linux file system hierarchy presents a collection of file systems on separate storage devices as if it were one set of files on one giant storage device that you can navigate.

File systems must be implemented in terms of raw disk blocks. Linux developed a well-defined kernel interface that allowed multiple types of file-systems to be active at once. A good file-system usually provides:

- Good performance
- Tolerance for crashes and power outages without file-system corruption
- The ability to handle disks and files large enough for your needs

Fortunately, Linux default file system cover all of these.

The “fourth extended file system, `ext4`, is the mainstream Linux standard. With journaling capability that increases reliability enormously. If a crash occurs, the file system uses the journal log to reconstruct a perfectly consistent file system. This also reduces the time needed to perform file system consistency checks to approximately one second per file system. `ext4` also raises size limits, increases the performance of certain operations, and allows the use of “extents” (disk block ranges) for storage allocation rather than just individual disk blocks.

Mount Points

To make the contents of a file system available in the file system hierarchy, it must be mounted on an empty directory. This directory is called a mount point. Once mounted, if you use `ls` to list that directory, you will see the contents of the mounted file system, and you can access and use those files normally. Many file systems are automatically mounted as part of the boot process.

Block Devices


Low-level access to storage devices in Linux is provided by a special type of file called a **block device**. These block devices must be formatted with a file system before they can be mounted.

Block device files are stored in the `/dev` directory, along with other device files. Device files are created automatically by the operating system. In Linux, the first SATA/PATA, SAS, SCSI, or USB hard drive detected is called `/dev/sda`, the second is `/dev/sdb`, and so on.

These names represent the entire hard drive. Other types of storage will have other forms of naming.

TYPE OF DEVICE	DEVICE NAMING PATTERN
SATA/SAS/USB-attached storage	<code>/dev/sd*</code>
virtio-blk paravirtualized storage	<code>/dev/vd*</code>
NVMe-attached storage (many SSDs)	<code>/dev/nvme*</code>
SD/MMC/eMMC storage (SD cards)	<code>/dev/mmcblk*</code>
Floppy disk drives.	<code>/dev/fd*</code>

TYPE OF DEVICE	DEVICE NAMING PATTERN
IDE Storage	/dev/hd*
Printers	/dev/lp*
Optical drives (CD/DVD readers and burners)	/dev/sr*
USB Drives	/dev/usb*
Virtual Console Terminal	/dev/vcs*
Terminals (Tele-Typewriter)	/dev/tty*
Device Mapper from LVM	/dev/dm*



Creating file systems

Let's say that we want to reformat a flash drive with a Linux native file system, rather than the FAT32 system they normally have by default. This involves two steps:

1. Create a new **partition** layout (optional).
2. Create a new **file system** on the drive.

Warning! In the following exercise, we are going to format a flash drive. Use a drive that contains nothing you care about because it will be erased! Again, make absolutely sure you are specifying the correct device name for your system, not the one shown in this example.

Identify the proper disk

Let's review the mounted disks in our system

```
1 [root@localhost ~]# fdisk -l | grep 'Disk /' | sort
2 Partition 1 does not start on physical sector boundary.
3 Disk /dev/loop0: 96.98 MiB, 101695488 bytes, 198624 sectors
4 Disk /dev/loop1: 8 KiB, 8192 bytes, 16 sectors
5 Disk /dev/mapper/fedora_localhost--live-home: 387.02 GiB,
415563251712 bytes, 811646976 sectors
6 Disk /dev/mapper/fedora_localhost--live-root: 70 GiB, 75161927680
bytes, 146800640 sectors
7 Disk /dev/mapper/fedora_localhost--live-swap: 7.74 GiB, 8308916224
bytes, 16228352 sectors
8 Disk /dev/sda: 167.68 GiB, 180045766656 bytes, 351651888 sectors
9 Disk /dev/sdb: 298.09 GiB, 320072933376 bytes, 625142448 sectors
10 Disk /dev/sdc: 5.46 TiB, 6001175125504 bytes, 11721045167 sectors
11 Disk /dev/sdd: 7.5 GiB, 8054112256 bytes, 15730688 sectors
12 Disk /dev/zram0: 4 GiB, 4294967296 bytes, 1048576 sectors
```

I used grep to reduce our output just enough to see the Disk names and sorted the output so I can easily identify which one is my inserted usb drive

I can by default ignore mapper drives because I know those represent LVM not USB.

I can also ignore loop because those are different devices again not the USB.

That leave me with just a few options:

```
1 Disk /dev/sda: 167.68 GiB, 180045766656 bytes, 351651888 sectors
2 Disk /dev/sdb: 298.09 GiB, 320072933376 bytes, 625142448 sectors
3 Disk /dev/sdc: 5.46 TiB, 6001175125504 bytes, 11721045167 sectors
4 Disk /dev/sdd: 7.5 GiB, 8054112256 bytes, 15730688 sectors
5 Disk /dev/zram0: 4 GiB, 4294967296 bytes, 1048576 sectors
```

I can ignore the ones that doesn't match the size of my USB drive and that leaves me with the proper answer

```
1 Disk /dev/sdd: 7.5 GiB, 8054112256 bytes, 15730688 sectors
```

I can double check with a couple commands.

First with `mount` that actually shows me the brand of my USB maxell

```
1 [root@localhost ~]# mount | grep sdd
2 /dev/sdd1 on /run/media/jmedinar/maxell type vfat
  (rw,nosuid,nodev,relatime,uid=1000,gid=1000,fmask=0022,dmask=0022,co
   depage=437,iocharset=ascii,shortname=mixed,showexec,utf8,flush,error
   s=remount-ro,uhelper=udisks2)
```

Or with the `lsblk` command

```
1 [root@localhost ~]# lsblk | grep sdd
2 sdd                         8:48   1  7.5G  0 disk
3 └─sdd1                      8:49   1  7.5G  0 part
  /run/media/jmedinar/maxell
```

Manipulating partitions with fdisk

The `fdisk` program allows us to interact directly with disk-like devices at a very low level. With this tool, we can edit, delete, and create partitions on the device. To work with our flash drive, we must first unmount it (if needed) and then invoke the `fdisk` program as follows:

```
1 [root@localhost ~]# umount /run/media/jmedinar/maxell
2 [root@localhost ~]# fdisk /dev/sdd
```

We will be presented with a series of screens and multiple options. Read carefully!

```
1 [root@localhost ~]# fdisk /dev/sdd
2
3 Welcome to fdisk (util-linux 2.36).
4 Changes will remain in memory only, until you decide to write them.
5 Be careful before using the write command.
6
7 Command (m for help):
```

We select the `m` option, notice that we must specify the device in terms of the entire device, not by partition number. :

```
1 Command (m for help): m
2
3 Help:
4
5   DOS (MBR)
6     a  toggle a bootable flag
7     b  edit nested BSD disklabel
8     c  toggle the dos compatibility flag
9
10  Generic
11    d  delete a partition
12    F  list free unpartitioned space
13    l  list known partition types
14    n  add a new partition
15    p  print the partition table
16    t  change a partition type
17    v  verify the partition table
18    i  print information about a partition
19
20  Misc
21    m  print this menu
22    u  change display/entry units
23    x  extra functionality (experts only)
24
25  Script
26    I  load disk layout from sfdisk script file
27    O  dump disk layout to sfdisk script file
28
29  Save & Exit
30    w  write table to disk and exit
31    q  quit without saving changes
32
33  Create a new label
34    g  create a new empty GPT partition table
35    G  create a new empty SGI (IRIX) partition table
36    o  create a new empty DOS partition table
37    s  create a new empty Sun partition table
```

The first thing we want to do is examine the existing partition layout. We do this by entering `p` to print the partition table for the device:

```

1 Command (m for help): p
2 Disk /dev/sdd: 7.5 GiB, 8054112256 bytes, 15730688 sectors
3 Disk model: USB Disk
4 Units: sectors of 1 * 512 = 512 bytes
5 Sector size (logical/physical): 512 bytes / 512 bytes
6 I/O size (minimum/optimal): 512 bytes / 512 bytes
7 Disklabel type: dos
8 Disk identifier: 0xb70b8e46
9
10 Device     Boot Start      End  Sectors  Size Id Type
11 /dev/sdd1        2320 15730687 15728368  7.5G  c W95 FAT32 (LBA)

```

In this example, we see a 7.5 GB device with a single partition /dev/sdd1. The partition is identified as a Windows 95 FAT32 partition. Some programs will use this identifier to limit the kinds of operations that can be done to the disk, but most of the time it is not critical to change it. We will change it to indicate a Linux partition.

We must first find out what ID is used to identify a Linux partition using the option l.

```

1 Command (m for help): l
2
3 00 Empty          24 NEC DOS          81 Minix / old Lin bf
Solaris
4 01 FAT12          27 Hidden NTFS Win  82 Linux swap / So c1
DRDOS/sec (FAT-
5 02 XENIX root    39 Plan 9          83 Linux           c4
DRDOS/sec (FAT-
6 03 XENIX usr     3c PartitionMagic  84 OS/2 hidden or  c6
DRDOS/sec (FAT-
7 04 FAT16 <32M    40 Venix 80286    85 Linux extended  c7
Syrinx
8 05 Extended       41 PPC PReP Boot   86 NTFS volume set da Non-
FS data
9 06 FAT16          42 SFS             87 NTFS volume set db CP/M
/ CTOS /
10 07 HPFS/NTFS/exFAT 4d QNX4.x      88 Linux plaintext de Dell
Utility
11 08 AIX            4e QNX4.x 2nd part 8e Linux LVM           df
BootIt
12 09 AIX bootable   4f QNX4.x 3rd part 93 Amoeba           e1 DOS
access
13 0a OS/2 Boot Manag 50 OnTrack DM    94 Amoeba BBT           e3 DOS
R/O
14 0b W95 FAT32      51 OnTrack DM6 Aux 9f BSD/OS           e4
SpeedStor
15 0c W95 FAT32 (LBA) 52 CP/M          a0 IBM Thinkpad hi ea
Linux extended
16 0e W95 FAT16 (LBA) 53 OnTrack DM6 Aux a5 FreeBSD           eb BeOS
fs
17 0f W95 Ext'd (LBA) 54 OnTrackDM6    a6 OpenBSD           ee GPT

```

```

18  10 OPUS           55 EZ-Drive        a7 NeXTSTEP      ef EFI
19  11 Hidden FAT12   56 Golden Bow     a8 Darwin UFS      f0
20  Linux/PA-RISC b
21  12 Compaq diagnost 5c Priam Edisk    a9 NetBSD       f1
    SpeedStor
22  14 Hidden FAT16 <3  61 SpeedStor      ab Darwin boot    f4
    SpeedStor
23  16 Hidden FAT16 secondary   63 GNU HURD or Sys af HFS / HFS+    f2 DOS
24  17 Hidden HPFS/NTF   64 Novell Netware  b7 BSDI fs        fb
    VMware VMFS
25  18 AST SmartSleep   65 Novell Netware  b8 BSDI swap      fc
    VMware VMKCORE
26  1b Hidden W95 FAT3  70 DiskSecure Mult  bb Boot Wizard hid fd
    Linux raid auto
27  1c Hidden W95 FAT3  75 PC/IX          bc Acronis FAT32 L fe
    LANstep
28
29  Aliases:
30    linux        - 83
31    swap         - 82
32    extended     - 05
33    uefi         - EF
34    raid         - FD
35    lvm          - 8E
36    linuxex      - 85
37

```

In the list we can see the 83 for a regular Linux partition or the option 8e for Linux LVM. To change it we use the option t.

```

1 | Command (m for help): t
2 | Selected partition 1
3 | Hex code or alias (type L to list all): 83
4 | Changed type of partition 'W95 FAT32 (LBA)' to 'Linux'.

```

This completes all the changes that we need to make. Up to this point, the device has **been untouched** (all the changes have been stored in memory, not on the physical device), so we will write the modified partition table to the device and exit. To do this, we enter w at the prompt.

```

1 | Command (m for help): w
2 | The partition table has been altered!

```

If we had decided to leave the device unaltered, we could have entered q at the prompt to quit.

Creating a file system with mkfs

With our partition editing done it's time to create a new file system on our flash drive. To do this, we will use `mkfs` command, that can create file systems in a variety of formats. To create an `ext4` file system on the device, we use the `-t` option to specify the type `ext4`, followed by the name of the device containing the partition we wish to format:

```
1 | [root@localhost ~]# mkfs -t ext4 /dev/sdd1
```

The program will display a lot of information when `ext4` is the chosen file system type. To re-format the device to its original `FAT32` file system, specify `vfat` as the file system type:

```
1 | [root@localhost ~]# mkfs -t vfat /dev/sdd1
```

This process of partitioning and formatting can be used anytime additional storage devices are added to the system.

Mounting

A file system residing on a removable storage device needs to be mounted in order to access it. The `mount` command allows the root user to manually mount a file system. The first argument of the `mount` command specifies the file system to mount. The second argument specifies the directory to use as the mount point in the file-system hierarchy.

There are two common ways to specify the file system on a disk partition to the `mount` command:

- With the name of the device file in `/dev` containing the file system.
- With the `UUID` written to the file system, a universally-unique identifier.

Mounting a device is relatively simple. You need to identify the device you want to mount, make sure the mount point exists, and mount the device on the mount point.

Identifying the Block Device

A hot-pluggable storage device, whether a hard disk drive `HDD` or solid-state device `SSD` in a server caddy, or a `USB` storage device, might be plugged into a different port each time they are attached to a system.

Use the `lsblk` command to list the details of a specified block device or all the available devices.

```
1 | [root@localhost ~]# lsblk
2 NAME           MAJ:MIN RM  SIZE RO TYPE
3 MOUNTPOINT
4 loop0          7:0    0   97M  1  loop
5 /var/lib/snapd/snap/core/9665
4 loop1          7:1    0     8K  1  loop
5 /var/lib/snapd/snap/openboard/14
5 sda            8:0    0 167.7G  0 disk
```

6	—sda1		8:1	0	1G	0	part	/boot
7	—sda2		8:2	0	166.7G	0	part	
8	—fedora_localhost--live-root	253:0		0	70G	0	lvm	/
9	—fedora_localhost--live-swap	253:1		0	7.7G	0	lvm	[SWAP]
10	—fedora_localhost--live-home	253:2		0	387G	0	lvm	/home
11	sdb		8:16	0	298.1G	0	disk	
12	—sdb1		8:17	0	298.1G	0	part	
13	—fedora_localhost--live-home	253:2		0	387G	0	lvm	/home
14	sdc		8:32	0	5.5T	0	disk	
15	—sdc1		8:33	0	128M	0	part	
16	—sdc2		8:34	0	5.5T	0	part	
	/run/media/jmedinar/Seagate Backup Plus Drive							
17	sdd		8:48	1	7.5G	0	disk	
18	—sdd1		8:49	1	7.5G	0	part	

If you know that you just added a 7.5G storage device with one partition, then you can guess from the preceding output that /dev/sdd1 is the partition that you want to mount.

Mounting by Block Device Name

The following example mounts the file system in the /dev/sdd1 partition on the directory /mnt/data.

```
1 | [root@localhost ~]# mount /dev/sdd1 /mnt/data
```

To mount a file system, the destination directory must already exist.

IMPORTANT

If the directory acting as mount point is not empty, any files copied to that directory before the file system was mounted are not accessible until the file system is unmounted again.

This approach works fine in the short run. However, the order in which the operating system detects disks can change if devices are added to or removed from the system. This will change the device name associated with that storage device. A better approach would be to mount by some characteristic built into the file system.

Mounting by File-system UUID

One stable identifier that is associated with a file system is its **UUID**, a very long hexadecimal number that acts as a universally-unique identifier. This **UUID** is part of the file system and remains the same as long as the file system is not recreated.

The **lsblk -fp** command lists the full path of the device, along with the **UUIDs** and mount points, as well as the type of file system in the partition. If the file system is not mounted, the mount point will be blank.

NAME	FSTYPE	FSVER	LABEL	UUID	FSAVAIL	FSUSE%	MOUNTPOINT
/dev/sda							
└─/dev/sda1	ext4	1.0	99dc15a-4811-43dc-b39c-c6812de80a83		636.7M	28%	/boot
└─/dev/sda2	LVM2_member	LVM2_001	XxDjmf-W1vl-q4fd-LUFK-Xg4I-a1Q-ChosLZ				
└─/dev/mapper/fedora_localhost--live-root	ext4	1.0	abf65797-144a-4588-ad65-17eec9b283f5		41.3G	35%	/
└─/dev/mapper/fedora_localhost--live-swap	swap	1	6dd9dea-310f-4216-a675-be18a0928f7e				[SWAP]
└─/dev/mapper/fedora_localhost--live-home	ext4	1.0	e7a61212-08c2-4dfa-bfc1-b39dbc51d50c		312.6G	13%	/home

Mount the file system by the UUID of the file system.

```
1 [root@localhost ~]# mount UUID="e7a61212-08c2-4dfa-bfc1-b39dbc51d50c" /mnt/data
```

Automatic mounting

If you are logged in and using the graphical desktop environment, it will automatically mount any removable storage media when it is inserted. The removable storage device is mounted at `/run/media/USERNAME/LABEL` where `LABEL` is an identifier, often the name given to the file system when it was created if one is available.

Before removing the device, you should unmount it manually.

If you are not using a removable device but a permanent one you have to include it into the `/dev/fstab` configuration to have automatically mounted at boot time.

The `fstab` is your operating system's file system table. The `fstab` is configured to look for specific file systems and mount them automatically in a desired way each and every time.

```
1 #
2 # /etc/fstab
3 # Created by anaconda on Thu Feb 27 19:31:06 2020
4 #
5 # Accessible filesystems, by reference, are maintained under
6 # '/dev/disk/'.
7 #
8 # After editing this file, run 'systemctl daemon-reload' to update
9 # systemd
10 #
11 /dev/mapper/fedora_localhost--live-root      /
12         ext4      defaults        1 1
13 UUID=99dc15a-4811-43dc-b39c-c6812de80a83   /boot
14         ext4      defaults        1 2
15 /dev/mapper/fedora_localhost--live-home       /home
16         ext4      defaults        1 2
17 /dev/mapper/fedora_localhost--live-swap        none
18         swap     defaults        0 0
```

The `fstab` uses the same mechanism as the `mount` command so we can mount by partition or `uuid`. Let's review the fields in the file:

```
1 | /dev/mapper/fedora_localhost--live-root      /
  ext4      defaults          1 1
```

1. device path or uid to mount
2. mount point
3. file system type
4. mounting options
5. dumping flag
6. passing flag

Mounting options

There's a large set of options available, but there's a handful or so of very common ones.

Option	Description
auto/noauto	Specify whether the partition should be automatically mounted on boot. You can block specific partitions from mounting at boot-up by using “noauto”.
exec/noexec	Specifies whether the partition can execute binaries. If you have a scratch partition that you compile on, then this would be useful, or maybe if you have /home on a separate file system. If you're concerned about security, change this to “noexec”.
ro/rw	ro is read-only, and rw is read-write. If you want to be able to write to a file-system as the user and not as root, you'll need to have rw specified.
sync/async	This one is interesting. “sync” forces writing to occur immediately on execution of the command, which is ideal for floppies and USB drives, but isn't entirely necessary for internal hard disks. What “async” does is allow the command to execute over an elapsed time period, perhaps when user activity dies down and the like. Ever get a message asking to your “wait while changes are being written to the drive?” This is usually why.
nouser/user	This allows the user to have mounting and unmounting privileges. An important note is that user automatically implies noexec so if you need to execute binaries and still mount as a user, be sure to explicitly use exec as an option.

These options are separated by a comma and no spaces, and can be put in any order. If you're not sure about the defaults, it's okay to explicitly state your options. Things that are mounted from temporary places (like USB) won't follow this basic pattern unless you created entries for them (by UUID) in fstab.

dumping

The next option is a binary value `0` for false and `1` for true for “dumping.” This is a pretty much out-dated method of backup for cases when the system went down. You should leave this as `0`.

passing

The last option is a numeric value for “passing.” This tells the system the order in which to perform a file system check `fsck`.

- `0` it will be skipped
- `1-n` it will be checked in that order

The root file system should always be `1` and other file systems can go afterward. This works best for journaling file systems like `ext3/4` and `ReiserFS`. Older file systems like `FAT16/32` and `ext2` can take a while, so it’s better to turn their `fscking` off and do it periodically yourself.

Unmounting

The shutdown and reboot procedures unmount all file systems automatically. As part of this process, any file system data cached in memory is flushed to the storage device thus ensuring that the file system suffers no data corruption.

To unmount a file system, the `umount` command expects the mount point as an argument.

```
1 | [root@localhost ~]# umount /mnt/data
```

Unmounting is not possible if the mounted file system is in use. For the `umount` command to succeed, all processes needs to stop accessing data under the mount point.

```
1 | [root@localhost ~]# cd /mnt/data
2 | [root@localhost data]# umount /mnt/data
3 | umount: /mnt/data: target is busy.
```

The `lsof` command lists all open files and the process accessing them in the provided directory. It is useful to identify which processes currently prevent the file system from successful unmount.

```
1 | [root@localhost data]# lsof /mnt/data
2 | COMMAND PID USER FD TYPE DEVICE SIZE/OFF NODE NAME
3 | bash    1593 root cwd DIR 253,17   6        128 /mnt/data
4 | lsof     2532 root cwd DIR 253,17  19        128 /mnt/data
5 | lsof     2533 root cwd DIR 253,17  19        128 /mnt/data
```

Once the processes are identified, an action can be taken, such as waiting for the process to complete or sending a `SIGTERM` or `SIGKILL` signal to the process.


```
1 | [root@host ~]# umount /mnt/data
```


Logical Volume Management

Logical Volume Management LVM creates a layer of abstraction over physical storage, allowing you to create logical storage volumes. With LVM in place, you are not bothered with physical disk sizes because the hardware is hidden from the software so it can be resized and moved without stopping applications or unmounting file systems. You can think of LVM as dynamic partitions.

Advantages:

- **Resize storage pools:** You can extend the logical space as well as reduce it without reformatting the disks.
- **Flexible storage capacity:** You can add more space by adding more disks and adding them to the pool of physical storage, thus you have flexible storage capacity.



Requesting new storage

Requesting storage in the server's world is normally something that happens outside of the server itself, you can as some might imagine open the server box and insert the new disk connecting it to cables and that will be considered as internal storage but that is very limiting to the number of slots in the server to connect storage and space not to say time consuming, risky and energy consuming in reality in most cases the request will be from some sort of external storage device [DAS, SAN, NAS, Cloud] that is already connected or accessible to the server.

As system administrator most of the time you won't be in charge of the storage itself so you should only care about the storage once it has been attached to the server and you can start using it. We refer most commonly to any type of storage as "LUN" Logical Unit Number, block device or simply "Disk".

Behind the scenes a lot of things might have happened you might be getting RAID storage, SATA, ATA, SCSI, SSD, it might have been connected over the network, serial or Fibre, it might be 10GbE connection or 1-GbE it might have redundancy or not, and high availability or not, etc etc...

But at the end all you care is about your LUN being visible to your Linux server!

Storage devices

Once the LUN is attached to your server we have to format it using `fdisk` command we add the new disk to be usable by LVM (8e).

First identify the disks available in the system and the one you will be adding

```
1 | [root@localhost ~]# fdisk -l | grep Disk
```

Double-check is not an already assigned disk

```
1 | [root@localhost ~]# lsblk  
2 | [root@localhost ~]# fdisk -c -u /dev/sdd
```

Start the `fdisk` program against the disk you want to format:

```
1 | [root@localhost ~]# fdisk /dev/sdd
```

Select the following options in the screens:


- n: add a new partition
- p: primary partition
- For partition number just hit Enter to select the default
- First Sector just hit Enter to select the default
- Last Sector just hit Enter to select the default
- t: to change the partition type
- Set 8e as code

Now let's review our new partition table

- p: to print the partition table

Finally, let's write the changes

- w: to write the changes



Done. We are done with all the physical activity! Let's get into the Logical LVM!

Physical Volumes

A physical volume is the actual storage device **representation** that will be used in the LVM configuration. It can be an entire disk or a partition of a disk.

PV Commands:

Command	Description
pvcreate	Initialize physical volume(s) for use by LVM
pvdisplay	Display various attributes of physical volume(s)
pvremove	Remove LVM label(s) from physical volume(s)
pvchange	Change attributes of physical volume(s)
pvmove	Move extents from one physical volume to another
pvresize	Resize physical volume(s)
pvscan	List all physical volumes
pvck	Check metadata on physical volumes
pvs	Display information about physical volumes

The **pvcreate** command initializes the disk so it can be a part of future volume group.


```
1 | [root@localhost ~]# pvcreate /dev/sdb
```

You can display PVs with the pvdisplay command

```
1 | [root@localhost ~]# pvdisplay
```

Volume Groups

Physical Volumes are combined into Volume Groups VGs. To create a pool of disk space out of which logical volumes can be allocated. The disk space available for allocation in Volume Group is divided into units of a fixed-size called extents. An extent is the smallest unit of storage that can be allocated. Within a physical volume, extents are referred to as physical extents PE. A logical volume is allocated into logical extents of the same size as the physical extents. The extent size is thus the same for all logical volumes in the volume group. The volume group maps the logical extents to physical extents.



So let's create a VG:

```
1 | [root@localhost ~]# vgcreate dataVG /dev/sdb
```

You can display information

```
1 | [root@localhost ~]# vgdisplay
```

Notice the VG Size and the PE Size those are important as well as the Total PE and Free PE Values.

VG Commands:

Command	Description
vgcreate	Create a volume group
vgdisplay	Display volume group information
vgremove	Remove volume group(s)
vgchange	Change volume group attributes

Command	Description
vgextend	Add physical volumes to a volume group
vgreduce	Remove physical volume(s) from a volume group
vgsplit	Move physical volumes into a new or existing volume group
vgmerge	Merge volume groups
vgexport	Unregister volume group(s) from the system
vgimport	Register exported volume group with system
vgrename	Rename a volume group
vgconvert	Change volume group metadata format
vgscan	Search for all volume group
vgck	Check the consistency of volume group(s)
vgs	Display information about volume groups
vgmknodes	Create the special files for volume group devices in /dev
vgimportclone	Import a VG from cloned PVs
vgcfgbackup	Backup volume group configuration(s)
vgcfgrestore	Restore volume group configuration

Logical Volumes

Once our volumes are grouped we will break them apart again but now to fit our needs into Logical Volumes. Logical Volumes (LV) are the final storage unit in the standard LVM architecture, and we have a few different options to cover unique use cases that we can employ to best fit a given situation.

Linear logical volume

Linear logical volumes are the LVM default when it comes to logical volume creation. They are generally used to combine one or more disks to create one usable storage unit.

Striped logical volume

Striped logical volumes allow the administrator to control the way that the data is written to the physical volumes. For high volume read/write scenarios, striped logical volumes would be ideal, as they allow for read and write operations to be done in parallel.

When using striped logical volumes, you can set the number of stripes (this number cannot exceed the number of physical volumes) and the stripe size. This allows the user a greater level of control over how I/O is performed on the system.

Mirrored logical volume

Mirrored logical volumes do exactly what you would expect them to do. They allow you to "reflect" the data on one device to an identical copy. This ensures your data is available. If one part of the mirror breaks, the remaining drive changes its characteristics to that of a linear volume and is still accessible.

Let's create a Linear logical volume

```
1 | [root@localhost ~]# lvcreate -L 1GB -n mylv myvg
```

You can display information

```
1 | [root@localhost ~]# lvdisplay
```

Notice the LV PAth, LV Status, and LV Size.

LV Commands:

Command	Description
lvcreate	Create a logical volume
lvdisplay	Display information about a logical volume
lvremove	Remove logical volume(s) from the system
lvchange	Change the attributes of logical volume(s)
lvextend	Add space to a logical volume
lvreduce	Reduce the size of a logical volume
lvresize	Resize a logical volume
lvrename	Rename a logical volume
lvconvert	Change logical volume layout
lvscan	List all logical volumes in all volume groups

Command	Description
lvs	Display information about logical volumes

File System

Create a file system from your LV as ext4 so it can be used by Linux

```
1 | [root@localhost ~]# mkfs -t ext4 /dev/vgname/lvname
```

Mounting

Mount the volume to any directory you need to

```
1 | [root@localhost ~]# mkdir /your/mount/point
2 | [root@localhost ~]# mount /dev/mapper/your_lvg /your/mount/point
```

You can also make the mount permanent so it will be available every time you reboot your server.

```
1 | [root@localhost ~]# vim /etc/fstab
```

add a line for your new mount as the following

```
1 | /dev/mapper/fedora_localhost--live-root / ext4 default 0 2
```

LVM partitioning is very useful if you have more than one hard disk, logical volumes can extend and are not limited by the size of one single disk, rather by the total aggregated size.


Managing Networking

When it comes to networking, there is probably nothing that cannot be done with Linux. Linux is used to build all sorts of networking systems and appliances, including firewalls, routers, name servers, NAS boxes and on and on. Just as the subject of networking is vast, so are the number of commands that can be used to configure and control it.

Networking concepts

TCP/IP Model

The TCP/IP network model is a simplified, four-layered set of abstractions that describes how different protocols interoperate in order for computers to send traffic from one machine to another over the Internet.



Network interface names

Each network port on a system has a name, which you use to configure and identify it.

Older versions of Linux used names like `eth0`, `eth1`, and `eth2` for each network interface. The name `eth0` was assigned first network port detected by the operating system, `eth1` the second, and so on. However, as devices are added and removed, the mechanism detecting devices and naming them could change which interface gets which name.

Newer versions of Linux use a different naming system. Instead of being based on detection order, the names of network interfaces are assigned based on information from the firmware, the PCI bus topology, and type of network device.

Naming system	Description
en	Ethernet Network
wl	[WLAN] Wireless Local Area Network Interface
ww	[WWAN] Wireless Wide Area Network Interface
oN	on-board device with N firmware index number
sN	PCI hotplug in slot N
pMsN	PCI on bus M in slot N
lo	loopback
virbr	Virtual bridge network
docker	docker bridge network

Examples:

```
1 eno1: Ethernet Network on-board 1
2 ens3: Ethernet Network PCI hotplug in slot 3
3 wlp4s0: Wireless PCI on bus 4 in slot 0
4 virbr0: Virtual Bridge 0
5 docker0: Docker Bridge 0
```

Persistent naming means that once you know what the name is for a network interface on the system, you also know that it will not change later. The trade off is that you cannot assume that a system with one interface will name that interface eth0.

IPv4

An IPv4 address is a 32-bit number, normally expressed in decimal as four 8-bit octets ranging in value from 0 to 255, separated by dots.

The address is divided into two parts:

- The network part (subnet).
- The host part (identifier).


All hosts on the same subnet can talk to each other directly without a router.

No two hosts on the same subnet can have the same host part.

The size of an IPv4 subnet is variable it depends on the netmask, which is assigned to the subnet. The netmask indicates how many bits of the IPv4 address belong to the subnet. The more bits available for the host part, the more hosts can be on the subnet.

The lowest possible address on a subnet is called the network address.

The highest possible address on a subnet is the broadcast address.



The special address 127.0.0.1 always points to the local system (“localhost”), and the network 127.0.0.0/8 belongs to the local system, so that it can talk to itself using network protocols.

IPv6

IPv6 is intended as an eventual replacement for the IPv4 network protocol. You will need to understand how it works since increasing numbers of production systems use IPv6 addressing.

IPv6 can also be used in parallel with IPv4 in a dual-stack model. In this configuration, a network interface can have an IPv6 address or addresses as well as IPv4 addresses. Linux operates in a dual-stack mode by default.

An IPv6 address is a 128-bit number, normally expressed as eight colon-separated groups of four hexadecimal nibbles. Each nibble represents four bits of the IPv6 address, so each group represents 16 bits of the IPv6 address.

```
1 | 2001:0db8:0000:0010:0000:0000:0000:0001
```

To make IPv6 addresses easier to write, leading zeros in a colon-separated group do not need to be written. However, at least one hexadecimal digit must be written in each colon-separated group.

```
1 | 2001:db8:0:10:0:0:0:1
```

Since addresses with long strings of zeros are common, one or more consecutive groups of zeros only may be combined with exactly one :: block.

```
1 | 2001:db8:0:10::1
```

Some tips for writing consistently readable addresses:

- Suppress leading zeros in a group.
- Use :: to shorten as much as possible.

- If an address contains two consecutive groups of zeros, equal in length, it is preferred to shorten the leftmost groups of zeros to :: and the rightmost groups to :0: for each group.
- Although it is allowed, do not use :: to shorten one group of zeros. Use :0: instead, and save :: for consecutive groups of zeros.
- Always use lowercase letters for hexadecimal numbers a through f.

Hostname

It would be inconvenient if you always had to use IP addresses to contact your servers. Humans generally would prefer to work with names than long and hard-to-remember strings of numbers. And so Linux has a number of mechanisms to map a host name to an IP address, collectively called name resolution.

One way is to set a static entry for each name in the /etc/hosts file on each system. This requires you to manually update each server's copy of the file.

Another way is using a DNS that is a distributed network of servers providing mappings of host names to IP addresses. In order for name service to work, a host needs to be pointed at a nameserver. This nameserver does not need to be on the same subnet; it just needs to be reachable by the host. This is typically configured through DHCP or a static setting in a file called /etc/resolv.conf.




The Internet: IP Addresses & DNS

Routing

Whether using IPv4 or IPv6, network traffic needs to move from host to host and network to network. Each host has a routing table, which tells it how to route traffic for particular networks. A routing table entry lists a destination network, which interface to use when sending traffic, and the IP address of any intermediate router required to relay a message to its final destination.

If the network traffic does not match a more specific route, the routing table usually has an entry for a default route to the entire IPv4 Internet: 0.0.0.0/0. This default route points to a router on a reachable subnet.

If a router receives traffic that is not addressed to it, instead of ignoring it like a normal host, it forwards the traffic based on its own routing table. This may send the traffic directly to the destination host, or it may be forwarded on to another router. This process of forwarding continues until the traffic reaches its final destination.



Gathering network information

Identifying Network Interfaces

The `ip link show` command will list all network interfaces available on your system:

```
1 [jmedinar@localhost ~]$ ip link show
2 1: lo: <LOOPBACK,UP,LOWER_UP> mtu 65536 qdisc noqueue state UNKNOWN
   mode DEFAULT group default qlen 1000
   link/loopback 00:00:00:00:00:00 brd 00:00:00:00:00:00
3 2: enp0s25: <NO-CARRIER,BROADCAST,MULTICAST,UP> mtu 1500 qdisc state DOWN
   mode DEFAULT group default qlen 1000
   link/ether 3c:97:0e:b7:3f:e4 brd ff:ff:ff:ff:ff:ff
4 3: wlp3s0: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc mq state UP
   mode DORMANT group default qlen 1000
   link/ether a4:4e:31:0d:90:50 brd ff:ff:ff:ff:ff:ff
```

To configure each network interface correctly, you need to know which one is connected to which network. In many cases, you will know the MAC address of the interface connected to each network, either because it is physically printed on the card or server, or because it is a virtual machine and you know how it is configured.

Displaying IP Addresses

Use the `ip addr show` command to view device and address information. A single network interface can have multiple IPv4 or IPv6 addresses.

```
1 [jmedinar@localhost ~]$ ip addr show wlp3s0
2 3: wlp3s0: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc mq state
   UP group default qlen 1000
3     link/ether a4:4e:31:0d:90:50 brd ff:ff:ff:ff:ff:ff
4       inet 192.168.1.39/24 brd 192.168.1.255 scope global dynamic
          noprefixroute wlp3s0
5         valid_lft 68201sec preferred_lft 68201sec
6       inet6 fe80::18ad:6abd:414f:60de/64 scope link noprefixroute
7         valid_lft forever preferred_lft forever
```

- An active interface state is UP.
- The link/ether line specifies the hardware (MAC) address of the device.
- The inet line shows an IPv4 address, its network prefix length, and scope.
- The inet6 line shows an IPv6 address, its network prefix length, and scope. This address is of global scope and is normally used.
- This inet6 line shows that the interface has an IPv6 address of link scope that can only be used for communication on the local Ethernet link.

Displaying Performance Statistics

The `ip -s link show` command may also be used to show statistics about network performance. Counters for each network interface can be used to identify the presence of network issues. The counters record statistics for things like the number of received (RX) and transmitted (TX) packets, packet errors, and packets that were dropped.

```
1 [jmedinar@localhost ~]$ ip -s link show wlp3s0
2 3: wlp3s0: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc mq state
   UP mode DORMANT group default qlen 1000
3     link/ether a4:4e:31:0d:90:50 brd ff:ff:ff:ff:ff:ff
4       RX: bytes packets errors dropped overrun mcast
5         1412387911 1265604    0        0        0        0
6       TX: bytes packets errors dropped carrier collsns
7         63637927   506098    0        0        0        0
```

Displaying the Routing Table

Use the `ip route` command with the `route` option to show routing information.

```
1 [jmedinar@localhost ~]$ ip route
2 default via 192.168.1.1 dev wlp3s0 proto dhcp metric 600
3 192.168.1.0/24 dev wlp3s0 proto kernel scope link src 192.168.1.39
metric 600
```

This shows the IPv4 routing table. All packets destined for the 192.168.1.0/24 network are sent directly to the destination through the device wlp3s0.

Add the -6 option to show the IPv6 routing table:

```
1 [jmedinar@localhost ~]$ ip -6 route
2 ::1 dev lo proto kernel metric 256 pref medium
3 fe80::/64 dev wlp3s0 proto kernel metric 600 pref medium
```

Displaying status of network devices

The nmcli dev status command displays the status of all network devices:

```
1 [jmedinar@localhost ~]$ nmcli dev status
2 DEVICE      TYPE      STATE           CONNECTION
3 wlp3s0      wifi      connected       loading-private
4 docker0     bridge    connected (externally) docker0
5 virbr0      bridge    connected (externally) virbr0
6 enp0s25    ethernet  unavailable    --
7 lo          loopback  unmanaged      --
8 virbr0-nic tun        unmanaged      --
```

The nmcli con show command displays a list of all connections. To list only the active connections, add the --active option.

```
1 [jmedinar@localhost ~]$ nmcli con show
2 NAME            UUID                                  TYPE
3 DEVICE
4 loading-private 587180df-01f7-4f98-9eeb-43e95bb58479  wifi
5 wlp3s0
6 docker0         0ae44952-464c-4640-8d30-b21c89194678  bridge
7 docker0
8 virbr0          643392aa-3c8e-4b24-a47d-2dbd42d705d5  bridge
9 virbr0
10 CougarWiFi     707cb698-bb56-4663-8121-2537adaa501e  wifi
11 --
12 loading...       a4404367-3c42-4363-930f-20c85c9d84b2  wifi
13 --
14 SpectrumSetup-7D 8358b974-a2ca-42e2-9382-a8c85d6dc85  wifi
15 --
16 Wired connection 1  feb4d793-7e49-398f-abfc-1711b93fdff8  ethernet
17 --
```

Network Configurations

The root user can make any necessary network configuration changes with nmcli.

However, regular users that are logged in on the local console can also make many network configuration changes to the system.

They have to log in at the system's keyboard to either a text-based virtual console or the graphical desktop environment to get this control. The logic behind this is that if someone is physically present at the computer's console, it's likely being used as a workstation or laptop and they may need to configure, activate, and deactivate wireless or wired network interfaces at will. By contrast, if the system is a server in the datacenter, generally the only users logging in locally to the machine itself should be administrators.

Regular users that log in using ssh do not have access to change network permissions without becoming root.

You can use the `nmcli gen permissions` command to see what your current permissions are.

```
1 [root@localhost ~]# nmcli gen permissions
2 PERMISSION
3 VALUE
4 org.freedesktop.NetworkManager.checkpoint-rollback
5 yes
6 org.freedesktop.NetworkManager.enable-disable-connectivity-check
7 yes
8 org.freedesktop.NetworkManager.enable-disable-network
9 yes
10 org.freedesktop.NetworkManager.enable-disable-statistics
11 yes
12 org.freedesktop.NetworkManager.enable-disable-wifi
13 yes
14 org.freedesktop.NetworkManager.enable-disable-wimax
15 yes
16 org.freedesktop.NetworkManager.enable-disable-wwan
17 yes
18 org.freedesktop.NetworkManager.network-control
19 yes
20 org.freedesktop.NetworkManager.reload
21 yes
22 org.freedesktop.NetworkManager.settings.modify.global-dns
23 yes
24 org.freedesktop.NetworkManager.settings.modify.hostname
25 yes
26 org.freedesktop.NetworkManager.settings.modify.own
27 yes
28 org.freedesktop.NetworkManager.settings.modify.system
29 yes
30 org.freedesktop.NetworkManager.sleep-wake
31 yes
32 org.freedesktop.NetworkManager.wifi.scan
33 yes
34 org.freedesktop.NetworkManager.wifi.share.open
35 yes
36 org.freedesktop.NetworkManager.wifi.share.protected
37 yes
```

Network Manager

NetworkManager is a daemon that monitors and manages network settings. In addition to the daemon, there is a GNOME Notification Area applet providing network status information. Command-line and graphical tools talk to NetworkManager and save configuration files in the /etc/sysconfig/network-scripts directory.

```
1 [jmedinar@localhost ~]$ ls -l /etc/sysconfig/network-scripts/
2 -rw-r--r--. 1 root root 366 Oct 11 15:11 ifcfg-my-wireless-
3 network...
4 -rw-----. 1 root root 29 Oct 11 15:11 keys-my-wireless-network...
```

Adding a network connection

The nmcli con add command is used to add new network connections.

The following command adds a new connection named eno2 for the interface eno2, which gets IPv4 networking information using DHCP and autoconnects on startup. It also gets IPv6 networking settings by listening for router advertisements on the local link. The name of the configuration file is based on the value of the con-name option, eno2, and is saved to the /etc/sysconfig/network-scripts/ifcfg-eno2 file.

```
1 [root@localhost ~]# nmcli con add con-name eno2 type ethernet ifname
2 eno2
3 Connection 'eno2' (5d978db4-8937-4aa5-bb1e-f6b9b5218cdf)
4 successfully added.
```

Controlling network connections

The nmcli con up name command activates the connection name on the network interface it is bound to. Note that the command takes the name of a connection, not the name of the network interface.

So let's check the names first

```
1 [root@localhost ~]# nmcli con show
2 NAME UUID TYPE
3 DEVICE
4 docker0 0ae44952-464c-4640-8d30-b21c89194678 bridge
5 docker0
6 virbr0 643392aa-3c8e-4b24-a47d-2dbd42d705d5 bridge
7 virbr0
8 CougarWiFi 707cb698-bb56-4663-8121-2537adaa501e wifi
9 ...
10 eno2 5d978db4-8937-4aa5-bb1e-f6b9b5218cdf ethernet
11 ...
12 Wired connection 1 feb4d793-7e49-398f-abfc-1711b93fdfff8 ethernet
13 ...
```

Now we can activate the connection

```
1 | [root@localhost ~]# nmcli con up static-eno2
```

The `nmcli dev disconnect` command disconnects the network interface device and brings it down.

```
1 | [root@host ~]# nmcli dev disconnect eno2
```

Modifying network connection settings

NetworkManager connections have two kinds of settings. There are static connection properties, configured by the administrator and stored in the configuration files in `/etc/sysconfig/network-scripts/ifcfg-*`. There may also be active connection data, which the connection gets from a DHCP server and which are not stored persistently.

To list the current settings for a connection, run the `nmcli con show name` command, where `name` is the name of the connection.

- Settings in lowercase are static properties that the administrator can change.
- Settings in all caps are active settings in temporary use for this instance of the connection.

```
1 | [root@localhost ~]# nmcli con show my-wireless-network
2 | connection.id:                      my-wireless-network
3 | connection.uuid:                    dasda-01f7-4f98-9eeb-
4 | connection.stable-id:                --
5 | connection.type:                   802-11-wireless
6 | connection.interface-name:         wlp3s0
7 | connection.autoconnect:            yes
8 | connection.autoconnect-priority:   0
9 | connection.autoconnect-retries:    -1 (default)
10 | connection.multi-connect:        0 (default)
11 | connection.auth-retries:          -1
12 | connection.timestamp:            1609272097
13 | ... output truncated ...
```

The `nmcli con mod name` command is used to change the settings for a connection. These changes are also saved in the `/etc/sysconfig/network-scripts/ifcfg-name` file for the connection.

Deleting a network connection

The `nmcli con del name` command deletes the connection named `name` from the system, disconnecting it from the device and removing the file `/etc/sysconfig/network-scripts/ifcfg-name`.

```
1 | [root@localhost ~]# nmcli con del eno2
2 | Connection 'eno2' (5d978db4-8937-4aa5-bb1e-f6b9b5218cdf)
   successfully deleted.
```

Manually configuring the network

It is also possible to configure the network by directly editing the connection configuration files.

Connection configuration files control the software interfaces for individual network devices. These files are usually named `/etc/sysconfig/network-scripts/ifcfg-name`, where name refers to the name of the device or connection that the configuration file controls. The following are standard variables found in the file used for static or dynamic IPv4 configuration.

STATIC	DYNAMIC	EITHER
BOOTPROTO=none	BOOTPROTO=dhcp	DEVICE=ens3
IPADDR0=172.25.250.10		NAME="static-ens3"
PREFIX0=24		ONBOOT=yes
GATEWAY0=172.25.250.254		UUID=f3e8(...)ad3e
DEFROUTE=yes		USERCTL=yes
DNS1=172.25.254.254		

In the static settings, variables for IP address, prefix, and gateway have a number at the end. This allows multiple sets of values to be assigned to the interface. The DNS variable also has a number used to specify the order of lookup when multiple servers are specified.

After modifying the configuration files, run `nmcli con reload` to make NetworkManager read the configuration changes. The interface still needs to be restarted for changes to take effect.

```
1 [root@localhost ~]# nmcli con reload
2 [root@localhost ~]# nmcli con down "static-ens3"
3 [root@localhost ~]# nmcli con up "static-ens3"
```

Changing the hostname

The `hostname` command displays or temporarily modifies the system's fully qualified host name.

```
1 [root@localhost ~]# hostname
2 localhost.localdomain
3 [root@localhost ~]# hostname temp.hostname
4 [root@localhost ~]# hostname
5 temp.hostname
```

A static host name may be specified in the `/etc/hostname` file.

```
1 [root@localhost ~]# cat /etc/hostname
2 localhost.localdomain
```

The `hostnamectl` command is used to modify this file and may be used to view the status of the system's fully qualified host name. If this file does not exist, the host name is set by a reverse DNS query once the interface has an IP address assigned.


```
1 [root@localhost ~]# hostnamectl
2     Static hostname: localhost.localdomain
3     Transient hostname: temp.hostname
4             Icon name: computer-laptop
5                 Chassis: laptop
6             Machine ID: be54ea963f8b493dac8eff001550f1ed
7                 Boot ID: d694981fb2cb4f6caecf44425248576e
8     Operating System: Fedora 33 (Workstation Edition)
9         CPE OS Name: cpe:/o:fedoraproject:fedora:33
10            Kernel: Linux 5.9.16-200.fc33.x86_64
11        Architecture: x86-64
```

Changing back our hostname

```
1 [root@localhost ~]# hostnamectl set-hostname localhost.localdomain
2 [root@localhost ~]# hostnamectl
3     Static hostname: localhost.localdomain
4     Transient hostname: fedora
5             Icon name: computer-laptop
6                 Chassis: laptop
7             Machine ID: be54ea963f8b493dac8eff001550f1ed
8                 Boot ID: d694981fb2cb4f6caecf44425248576e
9     Operating System: Fedora 33 (Workstation Edition)
10        CPE OS Name: cpe:/o:fedoraproject:fedora:33
11            Kernel: Linux 5.9.16-200.fc33.x86_64
12        Architecture: x86-64
13 [root@localhost ~]# hostname
14 localhost.localdomain
```

Configuring name resolution

Name resolution order



The stub resolver is used to convert host names to IP addresses or the reverse. It determines where to look based on the configuration of the `/etc/nsswitch.conf` file. By default, the contents of the `/etc/hosts` file are checked first.

```
1 [root@localhost ~]# cat /etc/hosts
2 127.0.0.1 localhost localhost.localdomain localhost4
   localhost4.localdomain4
3 ::1           localhost localhost.localdomain localhost6
   localhost6.localdomain6
```

If an entry is not found in the `/etc/hosts` file, by default the stub resolver tries to look up the hostname by using a DNS nameserver.

The `/etc/resolv.conf` file controls how this query is performed:

- search: a list of domain names to try with a short host name. Both this and domain should not be set in the same file; if they are, the last instance wins.
- nameserver: the IP address of a nameserver to query. Up to three nameserver directives may be given to provide backups if one is down.

```
1 [root@localhost ~]# cat /etc/resolv.conf
2 # Generated by NetworkManager
3 domain example.com
4 search example.com
5 nameserver 172.25.254.254
```

NetworkManager updates the `/etc/resolv.conf` file using DNS settings in the connection configuration files.

```
1 [root@localhost ~]# nmcli con mod ID ipv4.dns IP
2 [root@localhost ~]# nmcli con down ID
3 [root@localhost ~]# nmcli con up ID
4 [root@localhost ~]# cat /etc/sysconfig/network-scripts/ifcfg-ID
5 ...output omitted...
6 DNS1=8.8.8.8
7 ...output omitted...
```

The default behavior of `nmcli con mod ID ipv4.dns IP` is to replace any previous DNS settings with the new IP list provided. A + or - symbol in front of the `ipv4.dns` argument adds or removes an individual entry.

```
1 [root@localhost ~]# nmcli con mod ID +ipv4.dns IP
```

Testing DNS Name Resolution

The `host HOSTNAME` command can be used to test DNS server connectivity.

```
1 [root@localhost ~]# host classroom.example.com
2 classroom.example.com has address 172.25.254.254
3 [root@localhost ~]# host 172.25.254.254
4 254.254.25.172.in-addr.arpa domain name pointer
   classroom.example.com.
```

Monitoring & Troubleshooting



Checking connectivity between hosts

The `ping` command is used to test connectivity. The `ping` command sends a special network packet called an ICMP ECHO_REQUEST to a specified host. The command continues sending packets every second until `Ctrl+C` is pressed unless options are given to limit the number of packets sent.

```

1 [jmedinar@localhost ~]$ ping -c3 collin.edu
2 PING collin.edu (192.231.40.28) 56(84) bytes of data.
3 64 bytes from 192.231.40.28 (192.231.40.28): icmp_seq=1 ttl=113
time=16.8 ms
4 64 bytes from 192.231.40.28 (192.231.40.28): icmp_seq=2 ttl=113
time=16.4 ms
5 64 bytes from 192.231.40.28 (192.231.40.28): icmp_seq=3 ttl=113
time=16.8 ms
6
7 --- collin.edu ping statistics ---
8 3 packets transmitted, 3 received, 0% packet loss, time 2003ms
9 rtt min/avg/max/mdev = 16.367/16.681/16.847/0.222 ms

```

The ping6 command is the IPv6 version of ping in Red Hat Enterprise Linux. It communicates over IPv6 and takes IPv6 addresses, but otherwise works like ping.

```

1 [user@host ~]$ ping6 2001:db8:0:1::1
2 PING 2001:db8:0:1::1(2001:db8:0:1::1) 56 data bytes
3 64 bytes from 2001:db8:0:1::1: icmp_seq=1 ttl=64 time=18.4 ms
4 64 bytes from 2001:db8:0:1::1: icmp_seq=2 ttl=64 time=0.178 ms
5 64 bytes from 2001:db8:0:1::1: icmp_seq=3 ttl=64 time=0.180 ms
6 ^C
7 --- 2001:db8:0:1::1 ping statistics ---
8 3 packets transmitted, 3 received, 0% packet loss, time 2001ms
9 rtt min/avg/max/mdev = 0.178/6.272/18.458/8.616 ms
10 [user@host ~]$

```

When the remote host is not reachable we will get the following error

```

1 [jmedinar@localhost ~]$ ping 192.168.1.66
2 PING 192.168.1.66 (192.168.1.66) 56(84) bytes of data.
3 From 192.168.1.39 icmp_seq=1 Destination Host Unreachable
4
5 --- 192.168.1.66 ping statistics ---
6 1 packets transmitted, 0 received, +1 errors, 100% packet loss, time
0ms
7
8 [jmedinar@localhost ~]$ ping6 2001:db8:0:1::1
9 ping6: connect: Network is unreachable

```

ping also helps you identify slow network responses.

Tracing routes

To trace the path that network traffic takes to reach a remote host through multiple routers, use either traceroute or tracepath. This can identify whether an issue is with one of your routers or an intermediate one. Both commands use UDP packets to trace a path by default; however, many networks block UDP and ICMP traffic. The traceroute command has options to trace the path with UDP (-T), ICMP (-I), or TCP (-T) packets. Typically, however, the traceroute command is not installed by default.

```

1 [jmedinar@localhost ~]$ tracepath collin.edu

```

```

2 1?: [LOCALHOST]                                pmtu 1500
3 1: _gateway                                     3.653ms
4 1: _gateway                                     2.165ms
5 2: 142.254.130.213                            10.658ms
6 3: no reply
7 4: agg24.plantxmp01r.texas.rr.com           17.739ms
8 5: agg27.crtntxjt01r.texas.rr.com           17.816ms
9 6: agg21.dllatxl301r.texas.rr.com           11.855ms
10 7: 66.109.1.216                             38.482ms
11 8: 209-18-43-77.dfw10.tbone.rr.com          20.487ms
asymm 9
12 9: no reply
13 10: ae2.cs1.dfw2.us.zip.zayo.com          14.790ms
asymm 15
14 11: ae1.mcs1.dfw2.us.eth.zayo.com          14.259ms
15 12: ae5.mpr1.dfw5.us.zip.zayo.com          12.609ms
16 13: 208.184.23.101                          14.014ms
17 ...output omitted...

```

Each line in the output of `tracepath` represents a router or hop that the packet passes through between the source and the final destination. Additional information is provided as available, including the round trip timing (RTT) and any changes in the maximum transmission unit (MTU) size. The `asymm` indication means traffic reached that router and returned from that router using different (asymmetric) routes. The routers shown are the ones used for outbound traffic, not the return traffic.

The `traceroute6` and `traceroute -6` commands are the equivalent to `tracepath` and `traceroute` for IPv6.

Similar result with `traceroute`

```

1 [jmedinar@localhost ~]$ traceroute collin.edu
2 traceroute to collin.edu (192.231.40.28), 30 hops max, 60 byte
packets
3 1 _gateway (192.168.1.1) 2.561 ms 2.625 ms 2.594 ms
4 2 142.254.130.213 (142.254.130.213) 15.507 ms 15.569 ms 15.556
ms
5 3 tge0-0-8.frsdtx1701h.texas.rr.com (70.125.219.93) 25.760 ms
25.790 ms 25.770 ms
6 4 agg24.plantxmp01r.texas.rr.com (24.175.49.227) 23.226 ms
23.304 ms 23.264 ms
7 5 agg27.crtntxjt01r.texas.rr.com (24.175.36.177) 18.265 ms
23.235 ms 18.114 ms
8 6 agg21.dllatxl301r.texas.rr.com (24.175.49.0) 18.195 ms 13.736
ms 13.631 ms
9 7 bu-ether24.dllstx976iw-bcr00.tbone.rr.com (66.109.6.52) 18.560
ms 13.377 ms 13.338 ms
10 8 66.109.5.121 (66.109.5.121) 13.108 ms 13.885 ms 11.967 ms
11 9 * * *
12 10 ae2.cs1.dfw2.us.zip.zayo.com (64.125.26.202) 17.992 ms 18.262
ms 18.282 ms
13 11 ae1.mcs1.dfw2.us.eth.zayo.com (64.125.28.125) 17.859 ms
17.852 ms 17.810 ms

```

```

14  12  ae5.mpr1.dfw5.us.zip.zayo.com (64.125.28.86)  17.738 ms  17.547
     ms  12.127 ms
15  13  208.184.23.101 (208.184.23.101)  12.100 ms  17.540 ms  17.151
     ms

```

In the output, we can see that connecting from our system to `collin.edu` requires traversing at least 12 routers. For routers that provided identifying information, we see their hostnames, IP addresses, and performance data, which includes three samples of round-trip time from the local system to the router. For routers that do not provide identifying information (because of router configuration, network congestion, firewalls, etc.), we see asterisks.

Checking ports and services

TCP services use sockets as end points for communication and are made up of an IP address, protocol, and port number. Services typically listen on standard ports while clients use a random available port. Well-known names for standard ports are listed in the `/etc/services` file.

The `ss` command is used to display socket statistics. The `ss` command is meant to replace the older tool `netstat`, part of the `net-tools` package, which may be more familiar to some system administrators but which is not always installed.

```

1 [jmedinar@localhost ~]$ ss -ta
2 State      Recv-Q      Send-Q      Local Address:Port      Peer
3 Address:Port      Process
4 LISTEN      0          4096      0.0.0.0:hostmon
5      0.0.0.0:*
6 LISTEN      0          32        192.168.122.1:domain
7      0.0.0.0:*
8 LISTEN      0          4096      127.0.0.53%lo:domain
9      0.0.0.0:*
10 LISTEN      0          128       127.0.0.1:ipp
11      0.0.0.0:*
12 ESTAB       0          0          192.168.1.39:40090
13      192.168.1.9:8008
14 ESTAB       0          0          192.168.1.39:50954
15      192.168.1.9:8009
16 ESTAB       0          0          192.168.1.39:56884
17      216.58.194.110:https
18 ESTAB       0          0          192.168.1.39:52702
19      192.168.1.8:8009
20 ESTAB       0          0          192.168.1.39:33460
21      31.13.93.53:https
22 ESTAB       0          0          192.168.1.39:54006
23      64.233.180.188:hpvroom
24 ESTAB       0          0          192.168.1.39:52294
25      23.4.41.235:https
26 ESTAB       0          0          192.168.1.39:56038
27      34.216.200.139:https
28 ESTAB       0          0          192.168.1.39:38356
29      13.249.29.43:https

```

16	LISTEN	0	4096	[:]:hostmon	[::]:*
17	LISTEN	0	128	[:1]:ipp	[::]:*

Options for ss and netstat

OPTION	DESCRIPTION
-n	Show numbers instead of names for interfaces and ports.
-t	Show TCP sockets.
-u	Show UDP sockets.
-l	Show only listening sockets.
-a	Show all (listening and established) sockets.
-p	Show the process using the sockets.
-A inet	Display active connections (but not listening sockets) for the inet address family. That is, ignore local UNIX domain sockets. For ss, both IPv4 and IPv6 connections are displayed. For netstat, only IPv4 connections are displayed. (netstat -A inet6 displays IPv6 connections, and netstat -46 displays IPv4 and IPv6 at the same time.)

netstat

The netstat program is used to examine various network settings and statistics. Through the use of its many options, we can look at a variety of features in our network.

```

1 [jmedinar@localhost ~]$ netstat -ie
2 Kernel Interface table
3 lo: flags=73<UP,LOOPBACK,RUNNING> mtu 65536
4     inet 127.0.0.1 netmask 255.0.0.0
5         inet6 ::1 prefixlen 128 scopeid 0x10<host>
6             loop txqueuelen 1000 (Local Loopback)
7             RX packets 247 bytes 24817 (24.2 KiB)
8             RX errors 0 dropped 0 overruns 0 frame 0
9             TX packets 247 bytes 24817 (24.2 KiB)
10            TX errors 0 dropped 0 overruns 0 carrier 0 collisions 0
11
12 wlp3s0: flags=4163<UP,BROADCAST,RUNNING,MULTICAST> mtu 1500
13     inet 192.168.1.39 netmask 255.255.255.0 broadcast
14         192.168.1.255
15             inet6 fe80::18ad:6abd:414f:60de prefixlen 64 scopeid
16             0x20<link>
17                 ether a4:4e:31:0d:90:50 txqueuelen 1000 (Ethernet)
18                 RX packets 1291758 bytes 1424192191 (1.3 GiB)
19                 RX errors 0 dropped 0 overruns 0 frame 0
20                 TX packets 517621 bytes 65418705 (62.3 MiB)
```

The netstat command has hundreds of options but is more and more being replaced by the ss command!

```

1 [jmedinar@localhost ~]$ netstat --help
2 usage: netstat [-vWnNccf] [<Af>] -r           netstat {-V|--
3   version|-h|--help}
4       netstat [-vWnNcaeol] [<Socket> ...]
5       netstat { [-vWnNac] -I[<Iface>] | [-veenNac] -i | [-cnNe]
6   -M | -s [-6tuw] } [delay]
7
8       -r, --route          display routing table
9       -I, --interfaces=<Iface> display interface table for
10      <Iface>
11      -i, --interfaces     display interface table
12      -g, --groups         display multicast group
13      memberships
14      -s, --statistics    display networking statistics
15      (like SNMP)
16      -M, --masquerade    display masqueraded connections
17
18      -v, --verbose        be verbose
19      -W, --wide           don't truncate IP addresses
20      -n, --numeric         don't resolve names
21      --numeric-hosts     don't resolve host names
22      --numeric-ports      don't resolve port names
23      --numeric-users      don't resolve user names
24      -N, --symbolic       resolve hardware names
25      -e, --extend          display other/more information
26      -p, --programs        display PID/Program name for
27      sockets
28      -o, --timers          display timers
29      -c, --continuous      continuous listing
30
31      -l, --listening      display listening server sockets
32      -a, --all              display all sockets (default:
33      connected)
34      -F, --fib              display Forwarding Information
35      Base (default)
36      -C, --cache            display routing cache instead of
37      FIB
38      -Z, --context          display SELinux security context
39      for sockets
40
41      <Socket>={-t|--tcp} {-u|--udp} {-U|--udplite} {-S|--sctp} {-w|--raw}
42                  {-x|--unix} --ax25 --ipx --netrom
43      <AF>=Use '-6|-4' or '-A <af>' or '--<af>'; default: inet

```



Secure Communications

Linux can be administered remotely over the network and of course to do so it has to be over a secure channel by using the SSH protocol implemented by the OpenSSH tool.

What is OpenSSH?

OpenSSH implements the Secure Shell or SSH protocol in Linux systems. The SSH protocol enables systems to communicate in an encrypted and secure fashion over an insecure network. You can use the `ssh` command to create a secure connection to a remote system, authenticate as a specific user, and get an interactive shell session on the remote system as that user. You may also use the `ssh` command to run an individual command on the remote system without running an interactive shell.

SSH is a client-server service so it must be running in both ends of the connection on port 22.

We can connect to a remote server `remotehost` using the same user name as the current local user.

```
1 [user01@localhost ~]$ ssh remotehost
2 user01@remotehost's password: *****
3 ...output omitted...
4 [user01@remotehost ~]$
```

You can the `exit` command to log out of the remote system.

```
1 [user01@remotehost ~]$ exit
2 logout
3 Connection to remotehost closed.
4 [user01@localhost ~]$
```

We can also connect the remote server `remotehost` using a different username.

```
1 [user01@localhost ~]$ ssh user02@remotehost
2 user02@remotehost's password: *****
3 ...output omitted...
4 [user02@remotehost ~]$
```


We can execute a command on the remote host as any user user without accessing the remote interactive shell.

```
1 [user01@localhost ~]$ ssh user02@remotehost hostname
2 user02@remotehost's password: *****
3 remotehost.lab.example.com
4 [user01@localhost ~]$
```

Notice that the preceding command displayed the output in the local system's terminal.

SSH client for Windows?

The most popular one is probably PuTTY by Simon Tatham and his team. The PuTTY program displays a terminal window and allows a Windows user to open an SSH session on a remote host. PuTTY is available at <http://www.chiark.greenend.org.uk/~sgtatham/putty/>



↗ [How SSH Works](#)

SSH Keys

Generating a key pair provides you with two long string of characters: a public and a private key. You can place the public key on any server, and then unlock it by connecting to it with a client that already has the private key. When the two matches, the system unlocks without the need for a password. You can increase security even more by protecting the private key with a passphrase.

Step One—Create the RSA Key Pair

```
1 | [jmedinar@localhost ~]$ ssh-keygen -t rsa
```

Step Two—Store the Keys and Passphrase

Once you have entered the Gen Key command, you will get a few more questions:

```
[jmedinar@localhost ~]$ ssh-keygen -t rsa
Generating public/private rsa key pair.
Enter file in which to save the key (/home/jmedinar/.ssh/id_rsa):
Created directory '/home/jmedinar/.ssh'.
Enter passphrase (empty for no passphrase):
Enter same passphrase again:
Your identification has been saved in /home/jmedinar/.ssh/id_rsa.
Your public key has been saved in /home/jmedinar/.ssh/id_rsa.pub.
The key fingerprint is:
SHA256:uJkNEB015FkwOcootMwifTs5EnHb3FLVEEdvZ4lz70k jmedinar@localhost.localdomain
The key's randomart image is:
+---[RSA 3072]---+
|   +o=o...+=o   |
| .. .=o=. ...o .|
| +..=.=++     oo+o|
| +++ =.+..    .+o.|
| o. o oo.S      o|
| . = *        .|
| . o+ .      E |
|                         |
+---[SHA256]---
```

You can select where to save the keys, usually, the default value is the best place so just hit enter.

You can set a password for your keys. This is the same password you will use to access this machine where you are generating the keys. You can also leave it empty!

The public key is now located in your home directory under `.ssh/id_rsa.pub`. The private key (identification) is now located in the same location as `.ssh/id_rsa`.

```
1 [jmedinar@localhost ~]$ tree .ssh/
2 .
3   +-- id_rsa
4     -- id_rsa.pub
```

Log in with an SSH private key

Now that you have an SSH key pair, you're ready to configure your systems so they can SSH, SCP or SFTP to your server using your private key. You have to provide a copy of your **PUBLIC KEY ONLY!!!** to the remote host where you want to connect! **NEVER THE PRIVATE KEY!!!**

```
[jmedinar@localhost ~]$ cd .ssh/
[jmedinar@localhost .ssh]$ cat id_rsa.pub
ssh-rsa AAAAB3NzaC1yc2EAAAQABAAQBgQDIlPUoy0xvkei5WkEpBiGStx5drQ0yKRnyoml6
h1tQW07c3I52WwmIo0wr9z7Hr3coRGSAJGhk1YyjfIIKdFyNXoyZkN7JSKUF6w213zi0hQNDnBA
w6xIpo42Wz0XtXVs0FYyULtxgXFxe0B6Uks32d8fRHTB12vgk6JdiFSy9ZZGjPj2yI0dLgR6Aqs
Y6MegkUpJxWsMMjNPS7c= jmedinar@localhost.localdomain
[jmedinar@localhost ~]$ ls
```

The content of this file is a strange long string that represents your encrypted key!. You have to copy this content into the `~/.ssh/authorized_keys` file in the remote system. If the file doesn't exist you will have to create it.

```
1 | $(umask 077 && touch ~/.ssh/authorized_keys)
2 | [jmedinar@localhost ~]$ chmod 700 ~/.ssh
3 | [jmedinar@localhost ~]$ chmod 600 ~/.ssh/authorized_keys
```

Next, edit the file `.ssh/authorized_keys` using your preferred editor. Copy and paste your `id_rsa.pub` file into the file.

You can now SSH or SFTP into your server using your private key. From the command line, you can use:

```
1 | [jmedinar@localhost ~]$ ssh SYSUSER@x.x.x.x
```

If you didn't create your key in the default location, you'll need to specify the location:

```
1 | [jmedinar@localhost ~]$ ssh -i ~/.ssh/custom_key_name
SYSUSER@x.x.x.x
```

If you're using a Windows SSH client, such as PuTTy, look in the [configuration instructions](#)

scp [Secure Copy]

The OpenSSH package also includes two programs that can make use of an SSH-encrypted tunnel to copy files across the network. The first, scp (secure copy) is used much like the familiar cp program to copy files. The most notable difference is that the source or destination pathnames may be preceded with the name of a remote host, followed by a colon character.

The format is ALWAYS!!!

```
1 | [jmedinar@localhost ~]$ scp source target
```

To copy a file from our local system to our remote system we can do:

```
1 | [jmedinar@localhost ~]$ scp my_local_file.txt
root@remote_server:/root
```

and to copy from remote to local

```
1 | [jmedinar@localhost ~]$ scp
root@remote_server:/root/my_remote_file.txt /tmp
```

sftp [Secure FTP]

Ia a secure implementation for the ftp (File Transfer Protocol) program but it uses an SSH encrypted tunnel. sftp has an important advantage over conventional ftp in that it does not require an FTP server to be running on the remote host. It only requires the SSH server. This means that any remote machine that can connect with the SSH client can also be used as an FTP-like server.

```
1 | [jmedinar@localhost ~]$ sftp remote_server
```



The SFTP protocol is supported by many of the graphical file managers found in Linux distributions. Like [Filezilla](#).

[Using FileZilla Secure FTP or SFTP Connection](#)

<http://www.chiark.greenend.org.uk/~sgtatham/putty/>

SSH Tunneling

Part of what happens when you establish a connection with a remote host via SSH is that an encrypted tunnel is created between the local and remote systems. Normally, this tunnel is used to allow commands typed at the local system to be transmitted safely to the remote system, and for the results to be transmitted safely back. In addition to this basic function, the SSH protocol allows most types of network traffic to be sent through the encrypted tunnel, creating a sort of VPN (Virtual Private Network) between the local and remote systems.



[SSH Tunneling](#)